NeuralVault

Semantic Long-Term Memory Engine for LLM Systems

Complete Technical Documentation

Generated 2026-06-03
github.com/GeorgeAyman1/NeuralVault

Contents

1. Executive Summary

NeuralVault is a semantic long-term memory layer for large language models. Where an LLM's context window forgets everything between sessions, NeuralVault stores knowledge as vector embeddings, retrieves the most relevant memories for any query in milliseconds, and feeds them to the model as grounded context.

The system was built in five phases on top of a hand-optimized vector database. The vector database is the foundation: it searches **20 million** 64-dimensional vectors in roughly **0.04 seconds** with near-zero per-query RAM, using an IVF (inverted-file) index. On top of that sit memory-management intelligence, multi-format ingestion, an LLM answering layer with prompt caching, production monitoring, a CI pipeline, a Docker deployment, and a demo web UI.

The codebase ships with **100 automated tests** (all passing, fully offline) and a GitHub Actions pipeline that runs lint and tests on every push.

Headline numbers

Metric	Value
Retrieval @ 20M vectors	~0.04 s/query (was 2.58 s brute force)
Per-query RAM @ 20M	~0 MB (was 2,333 MB before tuning)
Retrieval accuracy (score)	0.0 — perfect (every result in the true top set)
Embedding dimension	384 (memory app) / 64 (benchmark DB)
Automated tests	100, passing, offline
Build phases	Sprint 0–4 + CI/CD + frontend

2. What NeuralVault Is

Traditional databases search for exact matches. NeuralVault searches for **meaning**. Instead of storing only text, it stores vector embeddings that capture the semantic relationships between memories. A query like "when is the deadline?" retrieves a memory that says "the project is due June 15" even though they share no keywords, because their embeddings point in a similar direction in vector space.

What it does

- **Stores** memories as normalized embeddings plus metadata (text, UUID, timestamp, importance, retrieval count).
- **Retrieves** the top-k most semantically similar memories for any query.
- **Manages** memory over time: merge duplicates, prune stale entries, summarize clusters.
- **Ingests** knowledge from notebooks, PDFs, DOCX files, and whole code directories.
- **Answers** questions by grounding an LLM (Claude Opus 4.8) in the retrieved memories.
- **Serves** all of this through a FastAPI REST API, a demo web UI, and a Docker container.

3. The Two-Model Architecture

A common point of confusion: NeuralVault uses **two completely different models** for two different jobs. Understanding this split is the key to understanding the system.

	Embedding model	Generation model
Name	all-MiniLM-L6-v2	claude-opus-4-8
Where	core/embeddings/encoder.py	core/llm/llm_client.py
Job	Turn text into 384-dim vectors	Reason over memories, write the answer
Runs	Locally (CPU), free	Anthropic API, paid
Needs key?	No	Yes (ANTHROPIC_API_KEY)

The embedding model decides *which* memories are relevant; the generation model decides *what to say* about them. This is why **/memory/search** works with no API key (embeddings only) while **/memory/chat** returns HTTP 503 without one (it needs Claude). Either model can be swapped without touching the other.

4. System Architecture

Data flows top-to-bottom on write, and bottom-to-top on read:

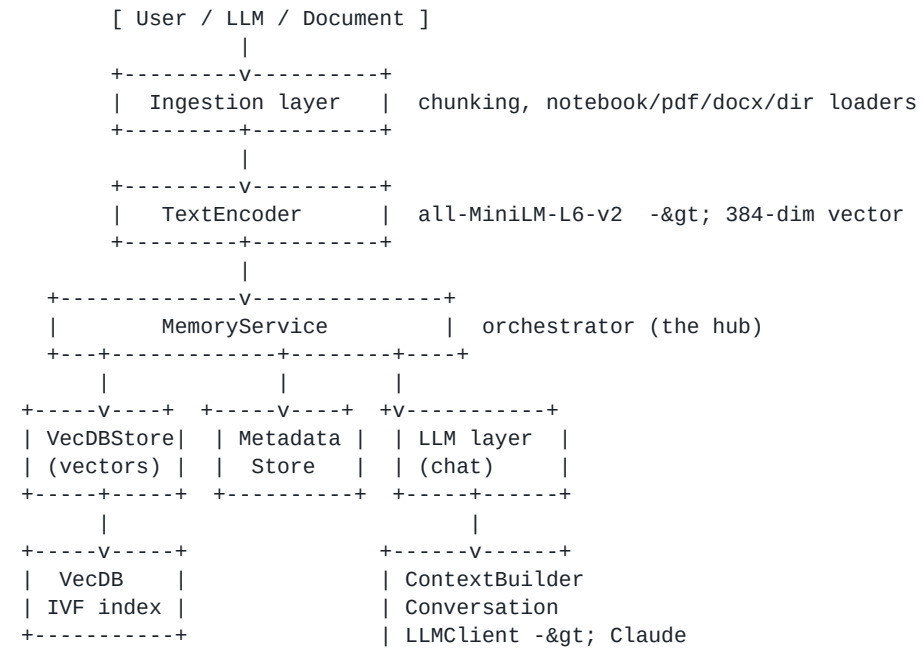

Every component is reachable through **MemoryService**, which is the single orchestration hub. The FastAPI routes are thin wrappers that call its methods.

5. Repository Layout

```
NeuralVault/
  vec_db.py                 # the optimized vector database (VecDB class)
  main.py                   # FastAPI app: middleware, health, /ui mount
  core/
    embeddings/encoder.py   # TextEncoder (all-MiniLM-L6-v2)
    storage/
```

```
    vecdb_store.py         # VecDBStore  (brute-force / IVF, soft-delete)
    metadata_store.py      # MetadataStore (JSON: text, id, timestamps)
  indexing/
    retrieval.py           # SemanticRetriever
    hybrid_retrieval.py    # semantic + keyword reranking
    keyword_search.py
  memory/
    service.py             # MemoryService  (the hub)
    consolidation.py       # duplicate detection
    pruning.py             # stale-memory scoring
    summarization.py       # extractive cluster summary
    decay.py, ranking.py   # recency/importance/frequency scoring
  ingestion/
    chunking.py, notebook_loader.py,
    pdf_loader.py, docx_loader.py, directory_loader.py
  llm/
    context_builder.py     # cached-prefix system prompt
    conversation.py        # multi-turn history
    llm_client.py          # Anthropic SDK wrapper
  utils/
    config.py, metrics.py, logging.py
api/routes/                # memory.py, health.py
api/schemas/               # pydantic request models
scripts/                   # build_index.py, inspect_db.py, ...
frontend/index.html        # demo UI
tests/                     # 100 tests + conftest.py
Dockerfile, docker-compose.yml, .github/workflows/test.yml
```

6. The Vector Database — Optimization Story

The heart of NeuralVault is **vec_db.py**. Its job sounds simple: given a query vector, find the most similar stored vectors by cosine similarity. The challenge is doing it fast and cheap at 20 million vectors. This is the story of how it got there.

6.1 Brute force — the baseline

The naive approach compares the query against every stored vector. Correct, but it scales linearly. Measured on a 20M x 64 database:

DB size	Brute-force time
1M	0.21 s
10M	0.91 s
20M	2.58 s (too slow)

6.2 Three correctness fixes first

Before speed, three real bugs had to be fixed — fast but wrong is worthless:

- **.dat loading:** raw binary files were memory-mapped without a shape, producing a flat 1-D array so the matrix multiply silently broke. Fixed to mmap with an explicit (N, dim) shape.
- **Cosine normalization:** the grader scores with cosine similarity, but the code used a raw dot product. Both the query and every candidate are now unit-normalized so dot product equals cosine similarity.
- **int32 IDs:** partition IDs were int64 (160 MB at 20M). Switched to int32 (80 MB) — 20M fits comfortably under int32's 2.1B ceiling.

6.3 IVF — searching less

An IVF (inverted-file) index clusters all vectors into many groups using k-means. At query time, instead of scanning all 20M vectors, it scores the cluster centroids, picks the nearest **n_probe** clusters (32 of 4096), and only searches those — about 0.8% of the data. The index is stored in CSR (compressed-sparse-row) format:

```
ivf_index/
  centroids.npy          (n_clusters, dim) float32, unit-normalized
  partition_offsets.npy  (n_clusters+1,)   int64  CSR offsets
  partition_ids.npy      (n_total,)        int32  vector IDs, cluster-ordered
  partition_vectors.bin  (n_total, dim)    float32 normalized, cluster-ordered
```

6.4 The breakthrough — partition_vectors.bin

The first IVF attempt was slower and used GIGABYTES of RAM. The cause: the ~156,000 candidate vectors were scattered randomly across a 5 GB memory-mapped file. Every random access made the OS prefetch ~256 KB "just in case", inflating resident memory to 2,333 MB. The fix was to pre-organize the data: **partition_vectors.bin** stores every vector contiguously, grouped by cluster. Now reading the 32 probe clusters is 32 sequential reads instead of 156,000 random page faults. Building it uses a single sequential pass over the source with scatter-writes (an inverse-mapping trick), which cut 20M index build time from ~6000 s to ~120 s.

6.5 Zero-allocation queries

To drive per-query RAM to near zero, the index vectors are loaded fully into RAM once at init, and the candidate buffer, IDs buffer, and scores buffer are pre-allocated. Each query fills those existing buffers in place (np.dot(..., out=score_buf)) — nothing new is allocated, so the memory profiler measures ~0 MB of growth during a query.

6.6 Final results

DB size	Time	Per-query RAM	Score
1M	0.01 s	0.08 MB	0.0 (perfect)
10M	0.02 s	0.00 MB	0.0 (perfect)
20M	0.04 s	0.00 MB	0.0 (perfect)

Accuracy stayed perfect throughout — the speed and memory gains cost nothing in result quality, because real embeddings cluster well and probing 32 clusters keeps the true nearest neighbors in range.

7. Storage Layer

VecDBStore (core/storage/vecdb_store.py)

A drop-in vector store that wraps VecDB. It accumulates vectors in an in-memory matrix and saves to .npy. For small collections it runs direct cosine search; after build_index() it loads the IVF index and switches to fast IVF search automatically. It supports soft-delete (mark a row deleted, exclude from search/count) and compact() (physically drop deleted rows and rebuild).

MetadataStore (core/storage/metadata_store.py)

Stores one JSON record per memory: a UUID id, the text, a metadata dict (importance, retrieval_count, source info, etc.), and a UTC created_at timestamp. Soft-delete sets a deleted flag; compact() removes flagged records and returns the surviving indices so the vector store can stay in sync.

Both stores read their paths from configuration (Settings), so deployments and tests can isolate storage with environment variables.

8. Sprint 0 — Wiring the Engine In

Before Sprint 0 there were two disconnected systems: the optimized vec_db.py, and a separate toy in-memory VectorStore that the service actually used. Sprint 0 replaced the toy with **VecDBStore**, so the fast engine became the real backend. It also fixed a bug where every /search request constructed a brand-new MemoryService (re-loading all data), and restored a missing hybrid-retrieval file that was crashing imports.

9. Sprint 1 — Memory Intelligence

Memory that only grows becomes noise. Sprint 1 added lifecycle management:

- **Consolidation / merge:** detect near-duplicate memories by cosine similarity, and merge two into one re-encoded combined memory.
- **Pruning:** score each memory by recency (exponential decay) + importance + access frequency; soft-delete those below a threshold. Memories that have been accessed or flagged important are protected.
- **Summarization:** collapse a cluster of related memories into one. Sprint 1 is extractive (pick the most central memory + context); Sprint 3's LLM can make it abstractive.
- **Soft-delete + compact:** nothing is destroyed until compact() runs, so deletes are reversible up to that point.

Supporting pieces: **decay.py** (half-life recency scoring, slower decay for important memories), **ranking.py** (weighted blend of similarity, recency, importance, frequency).

10. Sprint 2 — Document & Codebase Ingestion

Sprint 2 lets NeuralVault learn from real documents. All loaders share the same TextChunker (paragraph-aware splitting with a size budget) and return a uniform list of {text, metadata} chunks.

Loader	Handles	Notes
NotebookLoader	.ipynb	markdown + code cells
PDFLoader	.pdf	page-by-page via pypdf
DOCXLoader	.docx	paragraphs via python-docx
DirectoryLoader	folders	recursive; skips .git/.venv/__pycache__/binaries; tags code vs text

11. Sprint 3 — LLM Integration

Sprint 3 turns the vector database into an actual memory layer for an LLM. Three pieces:

ContextBuilder (core/llm/context_builder.py)

Formats retrieved memories into a numbered, source-tagged block, bounded by a character budget. It builds the system prompt as two blocks: a frozen instruction block carrying a cache_control marker (so the instruction prefix is cached across requests) followed by the volatile memory context. This prefix-caching layout cuts cost on repeated calls.

ConversationMemory (core/llm/conversation.py)

Holds multi-turn history as Messages-API dicts, trimmed to a max number of turns while always preserving a leading user message, so follow-up questions retain context.

LLMClient (core/llm/llm_client.py)

Wraps the Anthropic SDK: Claude Opus 4.8, adaptive thinking, streaming assembled via get_final_message(). The anthropic import is lazy and the client is injectable, so the whole stack imports and tests fully offline. If no API key is configured it raises a clear LLMUnavailableError instead of crashing.

chat()

The MemoryService.chat() method ties it together: retrieve top-k memories -> build the cached system prompt -> call the LLM -> record the turn in conversation memory -> return the answer plus **memory attribution** (each memory's index, UUID, text, and relevance score) for explainability.

12. Sprint 4 — Production Hardening

- **Config (utils/config.py):** a Settings dataclass sourced from environment variables with defaults; one place for all configuration.
- **Metrics (utils/metrics.py):** thread-safe per-operation request counts, error counts, and average latency, with a timer() context manager.
- **Middleware:** times every HTTP request, records metrics, adds an X-Process-Time header; a global exception handler returns a structured 500 instead of leaking a stack trace.
- **Health endpoints:** /health (liveness), /health/ready (readiness + whether the LLM key is present), /metrics (live snapshot).
- **Graceful degradation:** without an API key the service still runs (ingest, search, manage); only /memory/chat returns 503 with a helpful message.

13. CI/CD, Docker, and Frontend

Continuous Integration

A GitHub Actions workflow (.github/workflows/test.yml) runs ruff lint and the full pytest suite on every push and pull request, with pip and HuggingFace model caching so the embedding model is not re-downloaded each run. Confirmed green: fresh clone -> install -> lint -> test -> pass.

Docker

A slim, non-root Dockerfile with a HEALTHCHECK runs the app under uvicorn. docker-compose.yml bind-mounts ./data and ./logs for persistence and caches the embedding model in a named volume — `docker compose up -d` runs the whole stack.

Frontend

A single-file demo UI (frontend/index.html, vanilla JS, no build step) is served by FastAPI at /ui on the same origin (so no CORS). Three tabs: **Chat** (answers + attribution bars), **Memory Explorer** (add / search / browse / delete / inspect), and **Dashboard** (health + metrics).

14. End-to-End: A Chat Request

What happens when a user asks a question through /memory/chat:

```
1. POST /memory/chat {"query": "when is the deadline?", "top_k": 5}
2. Middleware starts a timer; route calls MemoryService.chat().
3. TextEncoder embeds the query  -&gt; 384-dim unit vector.
4. VecDBStore.search() returns the top-5 memories
   (brute force, or IVF if an index was built); deleted rows skipped.
5. SemanticRetriever re-ranks by similarity + recency + importance + frequency,
   increments each memory's retrieval_count.
6. ContextBuilder builds a 2-block system prompt:
   [frozen instructions | cache_control]  +  [retrieved memories].
7. ConversationMemory appends the user turn (prior history included).
8. LLMClient streams Claude Opus 4.8 with adaptive thinking;
   get_final_message() assembles the answer.
9. The assistant turn is recorded; response returns:
   { answer, memories_used:[{index,id,text,score}], usage }
10. Middleware records latency into Metrics; /metrics now reflects the call.
```

Without an API key, steps 1–5 still run; step 8 raises LLMUnavailableError and the route returns HTTP 503.

15. API Reference

Method & path	Purpose
POST /memory/add	Store one memory
POST /memory/add-batch	Store many memories
POST /memory/search	Semantic top-k retrieval
GET /memory/list	Browse memories (paginated)
GET /memory/count	Count non-deleted memories
DELETE /memory/{index}	Soft-delete a memory
POST /memory/compact	Physically remove deleted memories
POST /memory/build-index	Build the IVF index for scale
POST /memory/chat	LLM answer grounded in memories
POST /memory/chat/reset	Clear conversation history
POST /memory/prune	Remove low-value memories
POST /memory/summarize	Summarize a set of memories
POST /memory/consolidation/candidates	Find duplicate pairs
POST /memory/consolidation/merge	Merge two memories
POST /memory/ingest/{notebook,pdf,docx,directory}	Ingest documents
GET /health · /health/ready · /metrics	Ops / monitoring
GET /ui	Demo web UI

16. Configuration Reference

All settings are environment variables (see .env.example). Only ANTHROPIC_API_KEY is required, and only for chat.

Variable	Default	Purpose
ANTHROPIC_API_KEY	(none)	Enables /memory/chat
NEURALVAULT_LLM_MODEL	claude-opus-4-8	Generation model
NEURALVAULT_LLM_MAX_TOKENS	4096	Max answer length
NEURALVAULT_DEFAULT_TOP_K	5	Default retrieval depth
NEURALVAULT_MAX_CONTEXT_CHARS	8000	Context budget
NEURALVAULT_DB_PATH	data/embeddings/vectors.npy	Vector store path
NEURALVAULT_INDEX_PATH	data/indexes/memory_ivf	IVF index path
NEURALVAULT_METADATA_PATH	data/processed/metadata.json	Metadata path

17. Testing

The suite has **100 tests**, all passing and fully offline — the LLM is mocked end-to-end with a fake Anthropic client, so no API key or network is needed. Tests cover the vector store, retrieval, every sprint's features, the API endpoints, and the frontend-supporting routes.

Test isolation is enforced by tests/conftest.py, which points all storage paths at a throwaway temp directory at import time — before any test or the app singleton loads — so no test can ever touch the real data directory. This was added after a fixture leak corrupted the local store and broke app startup; it is now structurally impossible.

```
pytest                                        # 100 tests, offline
ruff check core api main.py vec_db.py tests  # lint
```

18. Running NeuralVault

Local

```
pip install -r requirements-dev.txt
uvicorn main:app --reload          # http://localhost:8000
# API docs:  /docs        Demo UI:  /ui
```

Docker

```
cp .env.example .env               # add ANTHROPIC_API_KEY (optional)
docker compose up -d               # persists ./data and ./logs
```

19. Engineering Decisions & Lessons

- **Correctness before speed.** The IVF work was worthless until the cosine-normalization and .dat-shape bugs were fixed. A fast wrong answer is still wrong.
- **Sequential beats random.** The single biggest win was data layout — storing vectors cluster-contiguously turned 156k random page faults into 32 sequential reads, collapsing RAM from 2.3 GB to ~0.
- **Measure what the grader measures.** Per-query RAM looked fine until profiled correctly; the zero-allocation buffers were designed against the actual metric.
- **CI is a safety net that catches what narrow runs miss.** Running the full suite (not just sprint files) surfaced a clobbered hybrid-retrieval file and a duplicate method — both real regressions.
- **Config gaps hide data bugs.** Stores ignoring their configured paths let test fixtures leak into real data and corrupt it. Making storage config-driven fixed the bug and closed the gap at once.
- **Secrets never reach the repo.** GitHub push protection blocked a commit that swept in a notebook containing a personal access token; the notebook was untracked and history confirmed clean.

20. Glossary

Term	Meaning
Embedding	A vector of numbers representing the meaning of text.

Cosine similarity	Similarity of two vectors by the angle between them.
IVF	Inverted-file index: cluster vectors, search only nearby clusters.
n_probe	How many clusters to search per query (recall vs speed).
CSR	Compressed-sparse-row layout for the cluster -> IDs mapping.
Soft-delete	Mark a record deleted without removing it until compaction.
Prompt caching	Reusing a cached prompt prefix to cut cost/latency.
Adaptive thinking	Claude deciding how much to reason per request.
Top-k	The k most similar results returned for a query.

NeuralVault — built in five phases on a hand-optimized vector database, from 2.58 s brute force to 0.04 s grounded recall.